

Using PostgreSQL

This guide focuses on the new PostgreSQL Service provided through [@cap-js/postgres](#), which is based on the same new database services architecture as the new [SQLite Service](#). This architecture brings significantly enhanced feature sets and feature parity, as documented in the [Features section of the SQLite guide](#).

↳ Learn about migrating from the former `cds-pg` in the [Migration chapter](#).

► *This guide is available for Node.js and Java.*

Table of Contents

- [Setup & Configuration](#)
 - [Auto-Wired Configuration](#)
- [Provisioning a DB Instance](#)
 - [Using Docker](#)
- [Service Bindings](#)
 - [Configure Service Bindings](#)
- [Deployment](#)
 - [Using cds deploy](#)
 - [With a Deployer App](#)
 - [Add PostgreSQL Deployment Configuration](#)
 - [Deploy](#)
- [Automatic Schema Evolution](#)
 - [Limitations](#)
 - [Dry-Run Offline](#)
 - [Generate Scripts](#)

- **Migration**
 - `cds deploy --model-only`
 - With Deployer App
- **MTX Support**

Setup & Configuration

Run this to use **PostgreSQL** for production:

```
npm add @cap-js/postgres
```

sh

Auto-Wired Configuration

The `@cap-js/postgres` package uses `cds-plugin` technique to auto-configure your application and use a PostgreSQL database for production.

You can inspect the effective configuration using `cds env`:

```
cds env requires.db --for production
```

sh

Output:

```
{  
  impl: '@cap-js/postgres',  
  dialect: 'postgres',  
  kind: 'postgres'  
}
```

js

↳ See also the general information on installing database packages

Provisioning a DB Instance

To connect to a PostgreSQL offering from the cloud provider in Production, leverage the [PostgreSQL on SAP BTP, hyperscaler option](#). For local development and testing convenience, you can run PostgreSQL in a [docker container](#).

Using Docker

You can use Docker to run a PostgreSQL database locally as follows:

1. Install and run [Docker Desktop](#)
2. Create the following file in your project root directory:

```
pg.yml  
-----  
services:  
  db:  
    image: postgres:alpine  
    environment: { POSTGRES_PASSWORD: postgres }  
    ports: [ '5432:5432' ]  
    restart: always  
yaml
```

3. Create and run the docker container:

```
docker-compose -f pg.yml up -d  
sh
```

Service Bindings

You need a service binding to connect to the PostgreSQL database.

In the cloud, use given techniques to bind a cloud-based instance of PostgreSQL to your application.

For local development provide the credentials using a suitable [cds env](#) technique, like one of the following.

Configure Service Bindings

Using Defaults with `[pg]` Profile

The `@cds-js/postgres` comes with default credentials under profile `[pg]` that match the defaults used in the [docker setup](#). So, in case you stick to these defaults you can skip the next sections and just go ahead, deploy your database:

```
cds deploy --profile pg
```

Run your application:

```
cds watch --profile pg
```

↳ Learn more about that in the [Deployment chapter below](#).

In Your private `~/.cdsrc.json`

Add it to your private `~/.cdsrc.json` if you want to use these credentials on your local machine only:

```
~/.cdsrc.json
```

```
{
  "requires": {
    "db": {
      "[pg)": {
        "kind": "postgres",
        "credentials": {
          "host": "localhost", "port": 5432,
          "user": "postgres",
          "password": "postgres",
          "database": "postgres"
        }
      }
    }
  }
}
```

In Project `.env` Files

Alternatively, use a `.env` file in your project's root folder if you want to share the same credentials with your team:

```
.env
```

```
cds.requires.db.[pg].kind = postgres                                properties
cds.requires.db.[pg].credentials.host = localhost
cds.requires.db.[pg].credentials.port = 5432
cds.requires.db.[pg].credentials.user = postgres
cds.requires.db.[pg].credentials.password = postgres
cds.requires.db.[pg].credentials.database = postgres
```

Using Profiles

The previous configuration examples use the `cds.env.profile [pg]` to allow selectively testing with PostgreSQL databases from the command line as follows:

```
cds watch --profile pg
```

sh

The profile name can be freely chosen, of course.

Deployment

Using `cds deploy`

Deploy your database as usual with that:

```
cds deploy
```

sh

Or with that if you used profile `[pg]` as introduced in the setup chapter above:

```
cds deploy --profile pg
```

sh

With a Deployer App

When deploying to Cloud Foundry, this can be accomplished by providing a simple deployer app. Similar to SAP HANA deployer apps, it is auto-generated for PostgreSQL-enabled projects by running

```
cds build --production
```

sh

- ▶ *What cds build does...*

Add PostgreSQL Deployment Configuration

```
cds add postgres
```

sh

- ▶ *See what this does...*

Deploy

You can package and deploy that application, for example using **MTA-based deployment**.

Automatic Schema Evolution

When redeploying after you changed your CDS models, like adding fields, automatic schema evolution is applied. Whenever you run `cds deploy` (or `cds-deploy`) it executes these steps:

1. Read a CSN of a former deployment from table `cds_model` .
2. Calculate the **delta** to current model.
3. Generate and run DDL statements with:
 - `CREATE TABLE` statements for new entities
 - `CREATE VIEW` statements for new views
 - `ALTER TABLE` statements for entities with new or changed elements

- *DROP & CREATE VIEW* statements for views affected by changed entities

4. Fill in initial data from provided .csv files using *UPSERT* commands.

5. Store a CSN representation of the current model in *cds_model*.

You can disable automatic schema evolution, if necessary, by setting *cds.requires.db.schema_evolution = false* *.

No manual altering

Manually altering the database will most likely break automatic schema evolution!

Limitations

Automatic schema evolution only allows changes without potential data loss.

Allowed

- Adding entities and elements
- Increasing the length of Strings
- Increasing the size of Integers

Disallowed

- Removing entities or elements
- Changes to primary keys
- All other type changes

For example the following type changes are allowed:

```
entity Foo {  
    anInteger : Int64;      // from former: Int32  
    aString : String(22);  // from former: String(11)  
}
```

cds

TIP

If you need to apply such disallowed changes during development, just drop and re-create your database, for example by killing it in docker and re-create it using the `docker-compose` command, [see Using Docker](#).

Dry-Run Offline

You can use `cds deploy` with option `--dry` to simulate and inspect how things work.

1. Capture your current model in a CSN file:

```
cds deploy --dry --model-only --out cds-model.csn
```

2. Change your models, for example in [capire/bookshop/db/schema.cds](#) :

```
entity Books { ...
    title : localized String(222); //> increase length from 111 to 222
    foo : Association to Foo;      //> add a new relationship
    bar : String;                 //> add a new element
}
entity Foo { key ID: UUID }      //> add a new entity
```

3. Generate delta DDL statements:

```
cds deploy --dry --delta-from cds-model.csn --out delta.sql
```

4. Inspect the generated SQL statements, which should look like this:

```
delta.sql
```

```
-- Drop Affected Views
DROP VIEW localized_CatalogService_ListOfBooks;
DROP VIEW localized_CatalogService_Books;
DROP VIEW localized_AdminService_Books;
DROP VIEW CatalogService_ListOfBooks;
DROP VIEW localized_sap_capire_bookshop_Books;
DROP VIEW CatalogService_Books_texts;
DROP VIEW AdminService_Books_texts;
DROP VIEW CatalogService_Books;
DROP VIEW AdminService_Books;
```

```

-- Alter Tables for New or Altered Columns
ALTER TABLE sap_capire_bookshop_Books ALTER title TYPE VARCHAR(222);
ALTER TABLE sap_capire_bookshop_Books_texts ALTER title TYPE VARCHAR(222);
ALTER TABLE sap_capire_bookshop_Books ADD foo_ID VARCHAR(36);
ALTER TABLE sap_capire_bookshop_Books ADD bar VARCHAR(255);

-- Create New Tables
CREATE TABLE sap_capire_bookshop_Foo (
    ID VARCHAR(36) NOT NULL,
    PRIMARY KEY(ID)
);

-- Re-Create Affected Views
CREATE VIEW AdminService_Books AS SELECT ... FROM sap_capire_bookshop_Boc
CREATE VIEW CatalogService_Books AS SELECT ... FROM sap_capire_bookshop_E
CREATE VIEW AdminService_Books_texts AS SELECT ... FROM sap_capire_books$t
CREATE VIEW CatalogService_Books_texts AS SELECT ... FROM sap_capire_book
CREATE VIEW localized_sap_capire_bookshop_Books AS SELECT ... FROM sap_c
CREATE VIEW CatalogService_ListOfBooks AS SELECT ... FROM CatalogService_
CREATE VIEW localized_AdminService_Books AS SELECT ... FROM localized_sap
CREATE VIEW localized_CatalogService_Books AS SELECT ... FROM localized_s
CREATE VIEW localized_CatalogService_ListOfBooks AS SELECT ... FROM local

```

Note: If you use SQLite, ALTER TYPE commands are not necessary and so, are not supported, as SQLite is essentially typeless.

Generate Scripts

You can use `cds deploy` with option `--script` to generate a script as a starting point for a manual migration. The effect of `--script` essentially is the same as for `--dry`, but it also allows changes that could lead to data loss and therefore are not supported in the automatic schema migration (see [Limitations](#)).

For generating such a script, perform the same steps as in section [Dry-Run Offline](#) above, but replace the command in step 3 by

```
cds deploy --script --delta-from cds-model.csn --out delta_script.sql
```

If your model change includes changes that could lead to data loss, there will be a warning and a respective comment is added to the dangerous statements in the resulting

script. For example, deleting an element or reducing the length of an element would look like this:

delta_script.sql

sql

```
...
-- [WARNING] this statement is lossy
ALTER TABLE sap_capire_bookshop_Books DROP price;

-- [WARNING] this statement could be lossy: length reduction of element "title"
ALTER TABLE sap_capire_bookshop_Books ALTER title TYPE VARCHAR(11);
...
```

WARNING

Always check and, if necessary, adapt the generated script before you apply it to your database!

Migration

Thanks to CAP's database-agnostic cds.ql API, we're confident that the new PostgreSQL service comes without breaking changes. Nevertheless, please check the instructions in the [SQLite Migration guide](#), which by and large applies also to the new PostgreSQL service.

`cds deploy --model-only`

Not a breaking change, but definitely required to migrate former `cds-pg` databases, is to prepare it for schema evolution.

To do so run `cds deploy` once with the `--model-only` flag:

`cds deploy --model-only`

sh

This will....:

- Create the `cds_model` table in your database.
- Fill it with the current model obtained through `cds compile '*'`.

IMPORTANT:

Your `.cds` models are expected to reflect the deployed state of your database.

With Deployer App

When you have a SaaS application, upgrade all your tenants using the [deployer app](#) with CLI option `--model-only` added to the start script command of your `package.json`. After having done that, don't forget to remove the `--model-only` option from the start script, to activate actual schema evolution.

MTX Support

WARNING

Multitenancy and extensibility aren't yet supported on PostgreSQL.

[Edit this page](#)

Last updated: 28/08/2025, 11:53

Previous page
[H2 \(Java\)](#)

Next page
[SAP HANA Cloud](#)

Was this page helpful?



