

Temporal Data

CAP provides out-of-the-box support for declaring and serving date-effective entities with application-controlled validity, in particular to serve as-of-now and time-travel queries.

Temporal data allows you to maintain information relating to past, present, and future application time. Built-in support for temporal data follows the general principle of CDS to capture intent with models while staying conceptual, concise, and comprehensive, and minimizing pollution by technical artifacts.

For an introduction to this topic, see [Temporal database](#) (Wikipedia) and [Temporal features in SQL:2011](#).

Table of Contents

- [Starting with 'Timeless' Models](#)
 - [Timeless Model](#)
 - [Timeless Data](#)
- [Declaring Temporal Entities](#)
 - [Using Annotations @cds.valid.from/to](#)
 - [Using Common Aspect temporal](#)
 - [Separate Temporal Details](#)
- [Serving Temporal Data](#)
- [Reading Temporal Data](#)
 - [As-of-now Queries](#)
 - [Time-Travel Queries](#)
 - [Time-Period Queries](#)
 - [Transitive Temporal Data](#)
- [Primary Keys of Time Slices](#)

Starting with 'Timeless' Models

For the following explanation, let's start with a base model to manage employees and their work assignments, which is free of any traces of temporal data management.

Timeless Model

timeless-model.cds

```
namespace com.acme.hr;                                         cds
using { com.acme.common.Persons } from './common';

entity Employees : Persons {
    jobs : Composition of many WorkAssignments on jobs.empl=$self;
    job1 : Association to one /*of*/ WorkAssignments;
}


entity WorkAssignments {
    key ID : UUID;
    role : String(111);
    empl : Association to Employees;
    dept : Association to Departments;
}

entity Departments {
    key ID : UUID;
    name : String(111);
    head : Association to Employees;
    members : Association to many Employees on members.jobs.dept = $self;
}
```

An employee can have several work assignments at the same time. Each work assignment links to one department.

Timeless Data

A set of sample data entries for this model, which only captures the latest state, can look like this:



Italic titles indicate to-one associations; actual names of the respective foreign key columns in SQL are `job1_ID`, `empl_ID`, and `dept_ID`.

Declaring Temporal Entities

Temporal Entities represent *logical* records of information for which we track changes over time by recording each change as individual *time slices* in the database with valid from/to boundaries. For example, we could track the changes of Alice's primary work assignment *WA1* over time:

Work Assignments					
empl	ID	validFrom	validTo	role	dept
E1	WA1	2014-01-01	2017-01-01	Developer	D2
E1	WA1	2017-01-01	2017-04-01	Developer	D1
E1	WA1	2017-04-01	2018-09-15	Senior Developer	D1
E1	WA1	2018-09-15	9999-12-31	Architect	D1
...

TIP

Validity periods are expected to be **non-overlapping** and **closed-open** intervals; same as in SQL:2011.

Using Annotations `@cds.valid.from/to`

To track temporal data, just add a pair of date/time elements to the respective entities annotated with `@cds.valid.from/to`, as follows:

```
entity WorkAssignments { //...  
    start : Date @cds.valid.from;  
    end   : Date @cds.valid.to;  
}
```

cds

TIP

The annotation pair `@cds.valid.from/to` actually triggers the built-in mechanisms for serving temporal data. It specifies which elements form the **application-time** period, similar to SQL:2011.

Using Common Aspect `temporal`

Alternatively, use the predefined aspect `temporal` to declare temporal entities:

```
using { temporal } from '@sap/cds/common';  
entity WorkAssignments : temporal {/*...*/}
```

cds

Aspect `temporal` is defined in `@sap/cds/common` as follows:

```
aspect temporal {  
    validFrom : Timestamp @cds.valid.from;  
    validTo   : Timestamp @cds.valid.to;  
}
```

cds

Separate Temporal Details

The previous samples would turn the whole `WorkAssignment` entity into a temporal one. Frequently though, only some parts of an entity are temporal, while others stay timeless. You can reflect this by separating temporal elements from non-temporal ones:

```
entity WorkAssignments {           // non-temporal head entity  
    key ID  : UUID;  
    empl  : Association to Employees;
```

cds


```

    details : Composition of WorkDetails on details.ID = $self.ID;
}

entity WorkDetails : temporal { // temporal details entity
    key ID : UUID; // logical record ID
    role : String(111);
    dept : Association to Departments;
}

```

The data situation would change as follows:



Serving Temporal Data

We expose the entities from the following timeless model in a service as follows:

service.cds

```

using { com.acme.hr } from './temporal-model';
service HRService {
    entity Employees as projection on hr.Employees;
    entity WorkAssignments as projection on hr.WorkAssignments;
    entity Departments as projection on hr.Departments;
}

```

cds

You can omit composed entities like `WorkAssignments` from the service, as they would get **auto-exposed** automatically.

Reading Temporal Data

As-of-now Queries

READ requests without specifying any temporal query parameter will automatically return data valid *as of now*.

For example, assumed the following OData query to read all employees with their current work assignments is processed on March 2019:

```
GET Employees? $expand=jobs($select=role&$expand=dept($select=name)) http
```

The values of `$at`, and so also the respective session variables, would be set to, for example:

<code>\$at.from = session_context('valid-from')</code>	<code>2019-03-08T22:11:00Z</code>
<code>\$at.to = session_context('valid-to')</code>	<code>2019-03-08T22:11:00.001Z</code>

The result set would be:

```
[ { "ID": "E1", "name": "Alice", "jobs": [ { "role": "Architect", "dept": { "name": "Core Development" } }, { "role": "Consultant", "dept": { "name": "App Development" } } ] }, { "ID": "E2", "name": "Bob", "jobs": [ { "role": "Builder", "dept": { "name": "Construction" } } ] }
```

Time-Travel Queries

We can run the same OData query as in the previous sample to read a snapshot data as valid on January 1, 2017 using the `sap-valid-at` query parameter:

```
GET Employees?sap-valid-at=date'2017-01-01'  
$expand=jobs($select=role&$expand=dept($select=name))
```

http

The values of `$at` and hence the respective session variables would be set to, for example:

```
$at.from = session_context('valid-from')= 2017-01-01T00:00:00Z  
$at.to = session_context('valid-to') = 2017-01-01T00:00:00.001Z
```

The result set would be:

```
[  
  { "ID": "E1", "name": "Alice", "jobs": [  
    { "role": "Developer", "dept": {"name": "Core Development"}},  
    { "role": "Consultant", "dept": {"name": "App Development"}}  
  ]}, ...  
]
```

WARNING

Time-travel queries aren't supported on SQLite due to the lack of `session_context` variables.

Time-Period Queries

We can run the same OData query as in the previous sample to read all history of data as valid since 2016 using the `sap-valid-from` query parameter:

```
GET Employees?sap-valid-from=date'2016-01-01'  
$expand=jobs($select=role&$expand=dept($select=name))
```

http

The result set would be:

```
{ "role": "Consultant", "dept": {"name": "App Development"} }  
], ...  
]
```

You would add `validFrom` in such time-period queries, for example:

```
GET Employees?sap-valid-from=date'2016-01-01'  
$expand=jobs($select=validFrom,role,dept/name)
```

http

WARNING

Time-series queries aren't supported on SQLite due to the lack of `session_context` variables.

TIP

Writing temporal data must be done in custom handlers.

Transitive Temporal Data

The basic techniques and built-in support for reading temporal data serves all possible use cases with respect to as-of-now and time-travel queries. Special care has to be taken though if time-period queries transitively expand across two or more temporal data entities.

As an example, assume that both, `WorkAssignments` and `Departments` are temporal:

```
using { temporal } from '@sap/cds/common';  
entity WorkAssignments : temporal /*...*/  
  dept : Association to Departments;  
}  
entity Departments : temporal /*...*/
```

cds

When reading employees with all history since 2016, for example:

```
GET Employees?sap-valid-from=date'2016-01-01'  
$expand=jobs(  
  $select=validFrom,role&$expand=dept(  
    $select=validFrom,name
```

http

```
)  
)
```

The results for *Alice* would be:

```
[  
  { "ID": "E1", "name": "Alice", "jobs": [  
    { "validFrom": "2014-01-01", "role": "Developer", "dept": [  
      {"validFrom": "2013-04-01", "name": "App Development"}  
    ]},  
    { "validFrom": "2017-01-01", "role": "Consultant", "dept": [  
      {"validFrom": "2013-04-01", "name": "App Development"}  
    ]},  
    { "validFrom": "2017-01-01", "role": "Developer", "dept": [  
      {"validFrom": "2014-01-01", "name": "Tech Platform Dev"},  
      {"validFrom": "2017-07-01", "name": "Core Development"}  
    ]},  
    { "validFrom": "2017-04-01", "role": "Senior Developer", "dept": [  
      {"validFrom": "2014-01-01", "name": "Tech Platform Dev"},  
      {"validFrom": "2017-07-01", "name": "Core Development"}  
    ]},  
    { "validFrom": "2018-09-15", "role": "Architect", "dept": [  
      {"validFrom": "2014-01-01", "name": "Tech Platform Dev"},  
      {"validFrom": "2017-07-01", "name": "Core Development"}  
    ]}  
  ], ...  
]
```

That is, all-time slices for changes to departments since 2016 are repeated for each time slice of work assignments in that time frame, which is a confusing and redundant piece of information. You can fix this by adding an alternative association to departments as follows:

```
using { temporal } from '@sap/cds/common';  
entity WorkAssignments : temporal /*...*/  
  dept : Association to Departments;  
  dept1 : Association to Departments on dept1.id = dept.id  
    and dept1.validFrom <= validFrom and validFrom < dept1.validTo;  
}  
entity Departments : temporal /*...*/
```

Primary Keys of Time Slices

While timeless entities are uniquely identified by the declared primary key — we call that the *conceptual* key in CDS — time slices are uniquely identified by *the conceptual key + validFrom*.

In effect the SQL DDL statement for the *WorkAssignments* would look like this:

```
CREATE TABLE com_acme_hr_WorkAssignments (                                sql
    ID : nvarchar(36),
    validFrom : timestamp,
    validTo : timestamp,
    --
    PRIMARY KEY ( ID, validFrom )
)
```

In contrast to that, the exposed API preserves the timeless view, to easily serve as-of-now and time-travel queries out of the box **as described above**:

```
<EntityType Name="WorkAssignments">                                xml
    <Key>
        <PropertyRef Name="ID"/>
    </Key>
    ...
</EntityType>
```

Reading an explicit time slice can look like this:

```
SELECT from WorkAssignments WHERE ID='WA1' and validFrom='2017-01-01'      sql
```

Similarly, referring to individual time slices by an association:

```
entity SomeSnapshotEntity {                                              cds
    //...
    workAssignment : Association to WorkAssignments { ID, validFrom }
}
```

[Previous page](#)[Localized Data](#)[Next page](#)[Security](#)

Was this page helpful?

